

RESEARCH MOTIVATION

I am captivated by the open challenges in machine learning that biology evolved to solve. Backpropagation can be inefficient due to global error signals while static architectures seem to lack the plasticity to prevent catastrophic forgetting. This rigidity in vision models appears to cause poor generalizability and vulnerability to adversarial attacks. Such fragilities in complex tasks can create intractable exploration spaces and unscalable communication. **I am motivated to research how integrating the principles of computation from biological substrates may reduce these limitations.**

SELECTED RESEARCH EXPERIENCES

My research with Drs. Hanna, Sala, and Berland has resulted in **eight first-author full-length papers: five peer-reviewed publications in NeurIPS workshops** (four in archival proceedings), one in NSF BRAID, and two preprints under review.

World Models for MARL. My research began with exploring bio-inspired multi-agent reinforcement learning (MARL), contributing to my team's **first-place finish at RoboCup**, the international MARL robotics competition. Leading a group investigating inter-agent communication, we found that an emergent strategy, inspired by signaling in simple organisms, collapsed to 12.2% success at task allocation with partial observability despite stabilization. Pivoting to experiment with a **world-model based approach** improved to 96.5% success, revealing the value of an agent internally simulating the environment. This taught me to remain open to unexpected results and shift focus based on data.

Active Inference & Structural Plasticity. Inspired by theories of intelligence like the free energy principle and Dr. Michael Levin's work on basal cognition, I addressed the intractability of **active inference** by developing a novel approximation using **principles from RL**. After grounding this work in discussions with Dr. Josiah Hanna (RL) and Dr. Levin (biology), **Dr. Levin invited me to extend this research under his mentorship**. Without a reward, the model maintained 82% success in Cart Pole, forming a step towards computable active inference. This research strengthened my ability to navigate the intersection of theoretical biology and machine learning.

Visual-Cortex Architecture. Curious about macro-level perspectives, I led a team experimenting with **primate visual cortex architectures** for light field identification. We implemented biological features including **dual-stream processing and predictive coding**. Despite engineering challenges of integration, the model achieved 74.4% accuracy, outperforming the next-best by 2.3 percentage points while being 2.5 times smaller, and demonstrated the value of inductive bias through geometric neuroscience. I learned to manage complex component interactions essential to building brain-inspired systems.

Industry Embodied AI. Concurrently, I lead **state estimation research** at an industry AI R&D lab. State-of-the-art (SOTA) algorithms proved unsuited for our constrained hardware and low-accuracy sensors. This drove me to develop an algorithm reducing dimensionality by disentangling data manifolds. The algorithm achieved over 100x improvement in accuracy over SOTA. It is now **deployed on all company robots** and received extensive positive customer feedback. This experience taught me to navigate the entire research-to-deployment pipeline under real-world constraints.

RESEARCH LEADERSHIP AND MENTORING

To create an interdisciplinary structure for this research at UW-Madison, **I founded and direct the Wisconsin Neuromorphic Computing and NeuroAI Lab**, securing formal funding, dedicated space, support from Dr. Akhilesh Jaiswal as advisor, and partnership with neuroAI startup FinalSpark. Through **nine graduate-level AI and neuroscience courses**, I gained technical foundations to lead this initiative. My role involves **mentoring 15 undergraduate researchers**, providing advice to over 100 researchers, developing research

proposals, organizing biweekly workshops, and lecturing on topics like spike-timing-dependent plasticity drawing audiences of over 100 undergraduates, graduates, and professionals.

SAMPLE RESEARCH EXTENSION

One fascinating direction building on my prior research investigates autonomous **emergence of state space** through predictive compression of high-dimensional sensory streams **during navigation**. While successor representations (SR) via spike timing-dependent plasticity are established, current models rely on pre-defined place cells or passive perception.

A **visual-cortex-based encoder** fuses with internal motion cues, and lateral inhibition encourages self-organization of sparse place cells. Distinctly, the model navigates to **minimize expected surprise** of its SR map. This research explores if such a bio-mimetic objective stabilizes the **online emergence of predictive maps**.

HOW A PHD FITS MY CAREER GOALS

Researching in industry, my development process was one of empirical iteration. I realized a deeper understanding of underlying mathematical theory may yield more efficient solutions. While I have strong practical skills, **I am driven to gain the theoretical depth to more rigorously devise novel algorithms**.

My long-term objective is conducting **research within a group like DeepMind's neuroscience lab**. Conversations with the current and previous lead, Drs. Kim Stachenfeld and Matthew Botvinick, solidified a PhD as the essential path to gain the theoretical depth and research freedom required. Receiving an **invitation from Dr. Karl Friston** to present my prospective doctoral research at his theoretical neurobiology group reinforced the value of a PhD for engaging with these ideas.

WHY MIT

I am applying to MIT because its unique concentration of faculty researching the application of neuroscience for AI provides the ideal environment to pursue my research goals, exemplified by the Quest for Intelligence, where I would love to contribute to the Quest for Developing Minds' work including reverse-engineering the intuitive physics in children.

Fascinated by **Dr. Joshua Tenenbaum's** work on human-like intelligence, especially "From Word Models to World Models", I aim to build on my experience in bio-inspired learning algorithms to avoid rejection sampling in Rational Meaning Construction by implementing a wake-sleep algorithm for amortized inference. High-fidelity posterior samples from the probabilistic engine would supervise the LLM, compiling symbolic reasoning steps into neural weights to transform the LLM into a proposal distribution.

Interested in **Dr. Nidhi Seethapathi's** work, especially "Exploration-based learning of a stabilizing controller predicts locomotor adaptation", I propose building on my experience in RL theory to extend the hierarchical framework to incorporate risk-sensitive RL replacing fixed noise with state-dependent noise. Scaling exploratory noise as a function of perceived stability, we could model how fear dynamically suppresses gradient estimation and investigate if reduced adaptation in impaired groups is an optimal policy response to physical instability.

Dr. James DiCarlo's research aligns with my experience computationally modeling the primate visual cortex. I could support the Integrative Research Program outlined in "How does the primate brain combine generative and discriminative computations in vision" through separating discriminative deepening and generative counter-streams by constructing a discriminative encoder refined by a generative objective via iterative amortized inference. Benchmarking against the temporal trajectory of inferior temporal cortex neural firing during ambiguous tasks would investigate if late-phase ventral stream dynamics exhibit normative signatures of probabilistic inference.

I am eager to bring my unique background in collaborative research and interdisciplinary curiosity to MIT and to contribute to its research community.