

RESEARCH MOTIVATION

I am captivated by the open challenges in machine learning that biology evolved to solve. Backpropagation can be inefficient due to global error signals while static architectures seem to lack the plasticity to prevent catastrophic forgetting. This rigidity in vision models appears to cause poor generalizability and vulnerability to adversarial attacks. Such fragilities in complex tasks can create intractable exploration spaces and unscalable communication. **I am motivated to research how integrating the principles of computation from biological substrates may reduce these limitations.**

SELECTED RESEARCH EXPERIENCES

My research with Drs. Hanna, Sala, and Berland has resulted in **eight first-author full-length papers: five peer-reviewed publications in NeurIPS workshops** (four in archival proceedings), one in BRAID, and two preprints under review.

World Models for MARL. My research began with exploring bio-inspired multi-agent reinforcement learning (MARL), contributing to my team's **first-place finish at RoboCup**, the international MARL robotics competition. Leading a group investigating inter-agent communication, we found that an emergent strategy, inspired by signaling in simple organisms, collapsed to 12.2% success at task allocation with partial observability despite stabilization. Pivoting to experiment with a **world-model based approach** improved to 96.5% success, revealing the value of an agent internally simulating the environment. This taught me to remain open to unexpected results and shift focus based on data.

Active Inference & Structural Plasticity. Inspired by theories of intelligence like the free energy principle and Dr. Michael Levin's work on basal cognition, I addressed the intractability of **active inference** by developing a novel approximation using **principles from RL**. After grounding this work in discussions with Dr. Josiah Hanna (RL) and Dr. Levin (biology), **Dr. Levin invited me to extend this research under his mentorship**. Without a reward, the model maintained 82% success in Cart Pole, forming a step towards computable active inference. This research strengthened my ability to navigate the intersection of theoretical biology and machine learning.

Visual-Cortex Architecture. Curious about macro-level perspectives, I led a team experimenting with **primate visual cortex architectures** for light field identification. We implemented biological features including **dual-stream processing and predictive coding**. Despite engineering challenges of integration, the model achieved 74.4% accuracy, outperforming the next-best by 2.3 percentage points while being 2.5 times smaller, and demonstrated the value of inductive bias through geometric neuroscience. I learned to manage complex component interactions essential to building brain-inspired systems.

Industry Embodied AI. Concurrently, I lead state estimation research at an industry AI R&D lab. State-of-the-art (SOTA) algorithms proved unsuited for our constrained hardware and low-accuracy sensors. This drove me to develop an algorithm reducing dimensionality by disentangling data manifolds. The algorithm achieved over 100x improvement in accuracy over SOTA. It is now **deployed on all company robots** and received extensive positive customer feedback. This experience taught me to navigate the entire research-to-deployment pipeline under real-world constraints.

RESEARCH LEADERSHIP AND MENTORING

To create an interdisciplinary structure for this research at UW-Madison, **I founded and direct the Wisconsin Neuromorphic Computing and NeuroAI Lab**, securing formal funding, dedicated space, support from Dr. Akhilesh Jaiswal as advisor, and partnership with neuroAI startup FinalSpark. Through **nine graduate-level AI and neuroscience**

courses, I gained technical foundations to lead this initiative. My role involves **mentoring 15 undergraduate researchers**, providing advice to over 100 researchers, developing research proposals, organizing biweekly workshops, and lecturing on topics like spike-timing-dependent plasticity drawing audiences of over 100 undergraduates, graduates, and professionals.

SAMPLE RESEARCH EXTENSION

One fascinating direction building on my prior research investigates autonomous **emergence of state space** through predictive compression of high-dimensional sensory streams **during navigation**. While successor representations (SR) via spike timing-dependent plasticity are established, current models rely on pre-defined place cells or passive perception.

A **visual-cortex-based encoder** fuses with internal motion cues, and lateral inhibition encourages self-organization of sparse place cells. Distinctly, the model navigates to **minimize expected surprise** of its SR map. This research explores if such a bio-mimetic objective **stabilizes the online emergence of predictive maps**.

HOW A PHD FITS MY CAREER GOALS

Researching in industry, my development process was one of empirical iteration. I realized a deeper understanding of underlying mathematical theory may yield more efficient solutions. While I have strong practical skills, **I am driven to gain the theoretical depth to more rigorously devise novel algorithms**.

My long-term objective is conducting **research within a group like DeepMind's neuroscience lab**. Conversations with the current and previous lead, Drs. Kim Stachenfeld and Matthew Botvinick, solidified a PhD as the essential path to gain the theoretical depth and research freedom required. Receiving an **invitation from Dr. Karl Friston** to present my prospective doctoral research at his theoretical neurobiology group reinforced the value of a PhD for engaging with these ideas.

WHY PRINCETON UNIVERSITY

Princeton University's unique faculty concentration in neuro-inspired AI provides the ideal environment for my research goals.

I am captivated by Dr. Tom Griffiths's resource-rational analysis. A potential research direction extends "Rational Metareasoning for LLMs" by drawing on my background in reinforcement learning. I propose investigating a Value of Computation (VOC) framework incorporating cost of external tool use. I aim to formulate policies where agents rationally weigh latency of API calls against internal parametric reasoning, investigating whether models autonomously minimize resource consumption while maintaining performance.

Dr. Brenden Lake's work on cognitive alignment is fascinating. While meta-learning for compositionality (MLC) achieves systematicity in "Human-Like Systematic Generalization Through a Meta-Learning Neural Network", it struggles with sequence lengths distinct from the training distribution. Leveraging my deep learning background, I propose extending MLC with recursive inductive biases or stack-augmented memory mechanisms to investigate length generalization bounds.

Dr. Sebastian Seung's approach to connectomics aligns with my research interests. Applying experience in computational neuroscience, I propose testing if biological wiring is structurally optimized for feature extraction. I aim to develop biologically-constrained CNNs, freezing kernels based on "Predicting Visual Function By Interpreting a Neuronal Wiring Diagram" Tm1-Dm3 and TmY connectivity matrices, to investigate if these structures yield superior data efficiency or robustness.

I am eager to contribute my background in collaborative research and interdisciplinary curiosity to Princeton's research community.

My research on neuroscience-inspired intelligent systems began with the scholarly challenge of academic isolation. Professors in AI told me that researching the application of neuroscience for AI is uncommon. Noting I was the only engineer in his class, a neuroscience professor expressed a wish for research on building AI's neurons into hardware, independently envisioning neuromorphic computing, unheard of in his department.

I had a passion for a field that, on my campus, did not seem to exist. To spare others this isolation, I created the community I once sought.

I founded the Wisconsin Neuromorphic Computing and NeuroAI Lab (WNCNL) to create both an interdisciplinary research hub and a learning community. I pitched the initiative to Dr. Akhilesh Jaiswal who offered his support as an advisor and I negotiated an official partnership with the NeuroAI startup FinalSpark. I was delighted to see the collaborative nature of the lab: students from neuroscience explaining the biomechanics of a neuron to engineers and computer scientists explaining artificial neural networks to biologists.

Leading this diverse group revealed the limits of my own foundational knowledge. I realized that to best contribute to this field, I require a deeper theoretical command of the relevant disciplines including mathematics, machine learning, and neuroscience, which drives my pursuit of graduate study.

I will bring this spirit of cross-pollination to Princeton. Having navigated the friction between disciplines, I am eager to serve as a mentor in the ReMatch and ReMatch+ programs, guiding undergraduates through the challenges of interdisciplinary research.