

RESEARCH MOTIVATION

I am captivated by the open challenges in machine learning that biology evolved to solve. Backpropagation can be inefficient due to global error signals while static architectures seem to lack the plasticity to prevent catastrophic forgetting. This rigidity in vision models appears to cause poor generalizability and vulnerability to adversarial attacks. Such fragilities in complex tasks can create intractable exploration spaces and unscalable communication. **I am motivated to research how integrating the principles of computation from biological substrates may reduce these limitations.**

SELECTED RESEARCH EXPERIENCES

My research with Drs. Hanna, Sala, and Berland has resulted in **eight first-author full-length papers: five peer-reviewed publications in NeurIPS workshops** (four in archival proceedings), one in NSF BRAID, and two preprints under review.

World Models for MARL. My research began with exploring bio-inspired multi-agent reinforcement learning (MARL), contributing to my team's **first-place finish at RoboCup**, the international MARL robotics competition. Leading a group investigating inter-agent communication, we found that an emergent strategy, inspired by signaling in simple organisms, collapsed to 12.2% success at task allocation with partial observability despite stabilization. Pivoting to experiment with a **world-model based approach** improved to 96.5% success, revealing the value of an agent internally simulating the environment. This taught me to remain open to unexpected results and shift focus based on data.

Active Inference & Structural Plasticity. Inspired by theories of intelligence like the free energy principle and Dr. Michael Levin's work on basal cognition, I addressed the intractability of **active inference** by developing a novel approximation using **principles from RL**. After grounding this work in discussions with Dr. Josiah Hanna (RL) and Dr. Levin (biology), **Dr. Levin invited me to extend this research under his mentorship**. Without a reward, the model maintained 82% success in Cart Pole, forming a step towards computable active inference. This research strengthened my ability to navigate the intersection of theoretical biology and machine learning.

Visual-Cortex Architecture. Curious about macro-level perspectives, I led a team experimenting with **primate visual cortex architectures** for light field identification. We implemented biological features including **dual-stream processing and predictive coding**. Despite engineering challenges of integration, the model achieved 74.4% accuracy, outperforming the next-best by 2.3 percentage points while being 2.5 times smaller, and demonstrated the value of inductive bias through geometric neuroscience. I learned to manage complex component interactions essential to building brain-inspired systems.

Industry Embodied AI. Concurrently, I **lead state estimation research** at an industry AI R&D lab. State-of-the-art (SOTA) algorithms proved unsuited for our constrained hardware and low-accuracy sensors. This drove me to develop an algorithm reducing dimensionality by disentangling data manifolds. The algorithm achieved over 100x improvement in accuracy over SOTA. It is now **deployed on all company robots** and received extensive positive customer feedback. This experience taught me to navigate the entire research-to-deployment pipeline under real-world constraints.

RESEARCH LEADERSHIP AND MENTORING

To create an interdisciplinary structure for this research at UW-Madison, **I founded and direct the Wisconsin Neuromorphic Computing and NeuroAI Lab**, securing formal funding, dedicated space, support from Dr. Akhilesh Jaiswal as advisor, and partnership with neuroAI startup FinalSpark. Through **nine graduate-level AI and neuroscience courses**, I gained technical foundations to lead this initiative. My role involves **mentoring 15 undergraduate researchers**, providing advice to over 100 researchers, developing research proposals, organizing biweekly workshops, and lecturing on topics like spike-timing-dependent plasticity drawing audiences of over 100 undergraduates, graduates, and professionals.

SAMPLE RESEARCH EXTENSION

One fascinating direction building on my prior research investigates autonomous **emergence of state space** through predictive compression of high-dimensional sensory streams **during**

navigation. While successor representations (SR) via spike timing-dependent plasticity are established, current models rely on pre-defined place cells or passive perception.

A **visual-cortex-based encoder** fuses with internal motion cues, and lateral inhibition encourages self-organization of sparse place cells. Distinctly, the model navigates to **minimize expected surprise** of its SR map. This research explores if such a bio-mimetic objective **stabilizes the online emergence of predictive maps**.

HOW A PHD FITS MY CAREER GOALS

Researching in industry, my development process was one of empirical iteration. I realized a deeper understanding of underlying mathematical theory may yield more efficient solutions. While I have strong practical skills, **I am driven to gain the theoretical depth to more rigorously devise novel algorithms.**

My long-term objective is conducting **research within a group like DeepMind's neuroscience lab**. Conversations with the current and previous lead, Drs. Kim Stachenfeld and Matthew Botvinick, solidified a PhD as the essential path to gain the theoretical depth and research freedom required. Receiving an **invitation from Dr. Karl Friston** to present my prospective doctoral research at his theoretical neurobiology group reinforced the value of a PhD for engaging with these ideas.

WHY UNIVERSITY COLLEGE LONDON

I am applying to University College London because its unique concentration of faculty researching applications of biology and neuroscience for AI provides the ideal environment to pursue my research goals.

Dr. **Peter Bentley**'s research on evolutionary developmental biology applied to AI is fascinating. "Development Necessitates Evolutionarily Conserved Factors" demonstrates that NCA can reproduce the disruption of anatomical structures upon perturbation. However, the current model fails to exhibit homeosis, likely due to the absence of metameric structures and reusable modules in the network architecture. Leveraging my background in deep learning and evolutionary computation, I propose to investigate whether true homeosis can be induced by re-engineering the NCA to utilize a hierarchical, modular controller rather than a flat, fully connected network. This research aims to rigorously determine if architectural modularity is a strict prerequisite for the emergence of homeotic developmental strategies, bridging continuous NCA representations and the discrete, modular nature of biological morphogenetic control.

I am captivated by Dr. **Tim Rocktäschel**'s pioneering work in open-endedness and unsupervised environment design. Building on the self-referential mechanisms introduced in "Promptbreeder", I propose investigating a co-evolutionary framework that surpasses the limits of static benchmarks by applying UED principles directly to the prompt space. While Promptbreeder effectively evolves solvers against fixed datasets like GSM8K, my research would explore a dual-population dynamic where an adversary evolves novel, progressively harder text-based tasks to generate an automated curriculum, while the solver population adapts via self-referential prompt evolution. Leveraging my background in deep reinforcement learning and evolutionary algorithms, I aim to rigorously analyze whether this adversarial interplay can prevent the plateauing often observed in static prompt optimization and unlock genuine open-ended reasoning capabilities in Large Language Models.

I am drawn to Dr. **David Barber**'s research on bayesian reasoning and probabilistic modeling. While "RotRNN" addresses the normalization and initialization instabilities of linear recurrent models by parameterizing the state matrix as a rotation, its current formulation as a linear time-invariant system limits its capacity for selective information flow compared to recent gated architectures like Mamba or Griffin. Drawing on my background in efficient deep learning architectures, I propose to extend this work by investigating an input-dependent gated rotational recurrence, where the rotation transition matrices become time-variant functions of the input, $A_t(u_t)$, to enable dynamic gating. My research would focus on deriving a parameterization of these dynamic rotations, potentially via time-varying block-diagonal structures or low-rank perturbations, that preserves the theoretical guarantees of $SO(N)$ orthogonality and constant norm while mathematically ensuring the recurrence remains compatible with parallel associative scans for efficient training.

I am eager to bring my unique background in collaborative research and interdisciplinary curiosity to University College London and to contribute to its research community.